

PADDING AND POOLING

- ① DNA can be expressed as a sequence of letters from the alphabet $\langle A, C, G, T \rangle$ ("adenine", "cytosine", "guanine", and "thymine"), e.g.

G A T T A C A A
C A G T T A A G
A A T A A A C C
G C A T T C A G

- ② Suppose that all DNA sequences of the form $*CA*TT*AG*$, where $*$ is an arbitrarily long (possibly empty) sequence, are responsible for some bodily function:


G A T T A C A A
C A G T T A A G ← positive
A A T A A A C C
G C A T T C A G ← positive

you're
it



PADDING AND POOLING

- ③ We can imagine building a CNN with three convolution kernels: one to detect "CA", one to detect "TT", and one to detect "AG".



- ④ The first thing to notice is that \dot{X} has one fewer column than X . This disparity will end up being rather annoying.

PADDING AND POOLING

⑤ We can remedy this by simply "padding" the end of the DNA sequence with a fifth letter, which we'll call "O".

$$\begin{array}{cccccccc}
 & C & A & G & T & T & A & A & G & (O) \\
 & \downarrow & \downarrow & \downarrow & \downarrow & \downarrow & \downarrow & \downarrow & \downarrow & \vdots \\
 X: & \begin{bmatrix} 0 & 1 & 0 & 0 & 0 & 0 & 1 & 1 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 0 \end{bmatrix} & \begin{matrix} 0 \\ 0 \\ 0 \\ 0 \end{matrix} \\
 & \begin{matrix} V & V & V & V & V & V & V & V & V \end{matrix} \\
 \dot{X}: & \begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \end{bmatrix} & \begin{matrix} \leftarrow "CA" \\ \leftarrow "TT" \\ \leftarrow "AG" \end{matrix}
 \end{array}$$

⑥ If we run this convolution layer on other positive sequences:

$$\begin{array}{ccc}
 GCATTTCAG & CATTTCAGC & TTCAATTAG \\
 VVVVVVVVV & VVVVVVVVV & VVVVVVVVV \\
 \dot{X}: \begin{bmatrix} 0 & 1 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \end{bmatrix} & \begin{bmatrix} 1 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \end{bmatrix} & \begin{bmatrix} 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \end{bmatrix}
 \end{array}$$

PADDING AND POOLING

⑦ Observe that the next layer of the CNN is still faced with a relatively daunting task, IF

$$\hat{X} = \begin{bmatrix} \hat{x}_{11} & \dots & \hat{x}_{18} \\ \hat{x}_{21} & \dots & \hat{x}_{28} \\ \hat{x}_{31} & \dots & \hat{x}_{38} \end{bmatrix}, \text{ then we want to predict}$$

a positive response in each of the following cases:

$$- \hat{x}_{11} = 1 \text{ and } \hat{x}_{23} = 1 \text{ and } \hat{x}_{35} = 1$$

$$- \hat{x}_{11} = 1 \text{ and } \hat{x}_{23} = 1 \text{ and } \hat{x}_{36} = 1$$

$$- \hat{x}_{11} = 1 \text{ and } \hat{x}_{23} = 1 \text{ and } \hat{x}_{37} = 1$$

$$- \hat{x}_{11} = 1 \text{ and } \hat{x}_{24} = 1 \text{ and } \hat{x}_{36} = 1$$

$$- \hat{x}_{11} = 1 \text{ and } \hat{x}_{24} = 1 \text{ and } \hat{x}_{37} = 1$$

$$- \hat{x}_{11} = 1 \text{ and } \hat{x}_{25} = 1 \text{ and } \hat{x}_{36} = 1$$

$$- \hat{x}_{12} = 1 \text{ and } \hat{x}_{24} = 1 \text{ and } \hat{x}_{36} = 1$$

$$- \hat{x}_{12} = 1 \text{ and } \hat{x}_{24} = 1 \text{ and } \hat{x}_{37} = 1$$

$$- \hat{x}_{12} = 1 \text{ and } \hat{x}_{25} = 1 \text{ and } \hat{x}_{37} = 1$$

$$- \hat{x}_{13} = 1 \text{ and } \hat{x}_{25} = 1 \text{ and } \hat{x}_{37} = 1$$


e.g., read this line as: "CA" starts at position 1, "TT" starts at position 3, and "AG" starts at position 7




There's no way for it to generalize between these cases, so the training data would need to contain several instances of each case.

PADDING AND POOLING

- ⑧ This only gets worse as the length of the DNA sequence grows. If its length is N , then the number of different cases to consider is cubic in N :



- ⑨ One way to tackle this issue is by creating a layer that summarizes the results of the detectors over subregions:



PADDING AND POOLING

⑩ This technique:

$$\begin{array}{c} \dot{X}: \begin{bmatrix} \dot{x}_{11} & \dot{x}_{12} & \dot{x}_{13} & \dot{x}_{14} & \dot{x}_{15} & \dot{x}_{16} & \dot{x}_{17} & \dot{x}_{18} \\ \dot{x}_{21} & \dot{x}_{22} & \dot{x}_{23} & \dot{x}_{24} & \dot{x}_{25} & \dot{x}_{26} & \dot{x}_{27} & \dot{x}_{28} \\ \dot{x}_{31} & \dot{x}_{32} & \dot{x}_{33} & \dot{x}_{34} & \dot{x}_{35} & \dot{x}_{36} & \dot{x}_{37} & \dot{x}_{38} \end{bmatrix} \\ \begin{array}{cccc} \vee & \vee & \vee & \vee \end{array} \\ \ddot{X}: \begin{bmatrix} \max(\dot{x}_{11}, \dot{x}_{12}) & \max(\dot{x}_{13}, \dot{x}_{14}) & \max(\dot{x}_{15}, \dot{x}_{16}) & \max(\dot{x}_{17}, \dot{x}_{18}) \\ \max(\dot{x}_{21}, \dot{x}_{22}) & \max(\dot{x}_{23}, \dot{x}_{24}) & \max(\dot{x}_{25}, \dot{x}_{26}) & \max(\dot{x}_{27}, \dot{x}_{28}) \\ \max(\dot{x}_{31}, \dot{x}_{32}) & \max(\dot{x}_{33}, \dot{x}_{34}) & \max(\dot{x}_{35}, \dot{x}_{36}) & \max(\dot{x}_{37}, \dot{x}_{38}) \end{bmatrix} \end{array}$$

is called max-pooling.

⑪ Notice that max-pooling greatly simplifies the next layer of the CNN. IF $\ddot{X} = \begin{bmatrix} \ddot{x}_{11} & \ddot{x}_{12} & \ddot{x}_{13} & \ddot{x}_{14} \\ \ddot{x}_{21} & \ddot{x}_{22} & \ddot{x}_{23} & \ddot{x}_{24} \\ \ddot{x}_{31} & \ddot{x}_{32} & \ddot{x}_{33} & \ddot{x}_{34} \end{bmatrix}$, then

We want to predict a positive response in these cases:

$$- \ddot{x}_{11} = 1 \text{ and } \ddot{x}_{22} = 1 \text{ and } \ddot{x}_{33} = 1$$

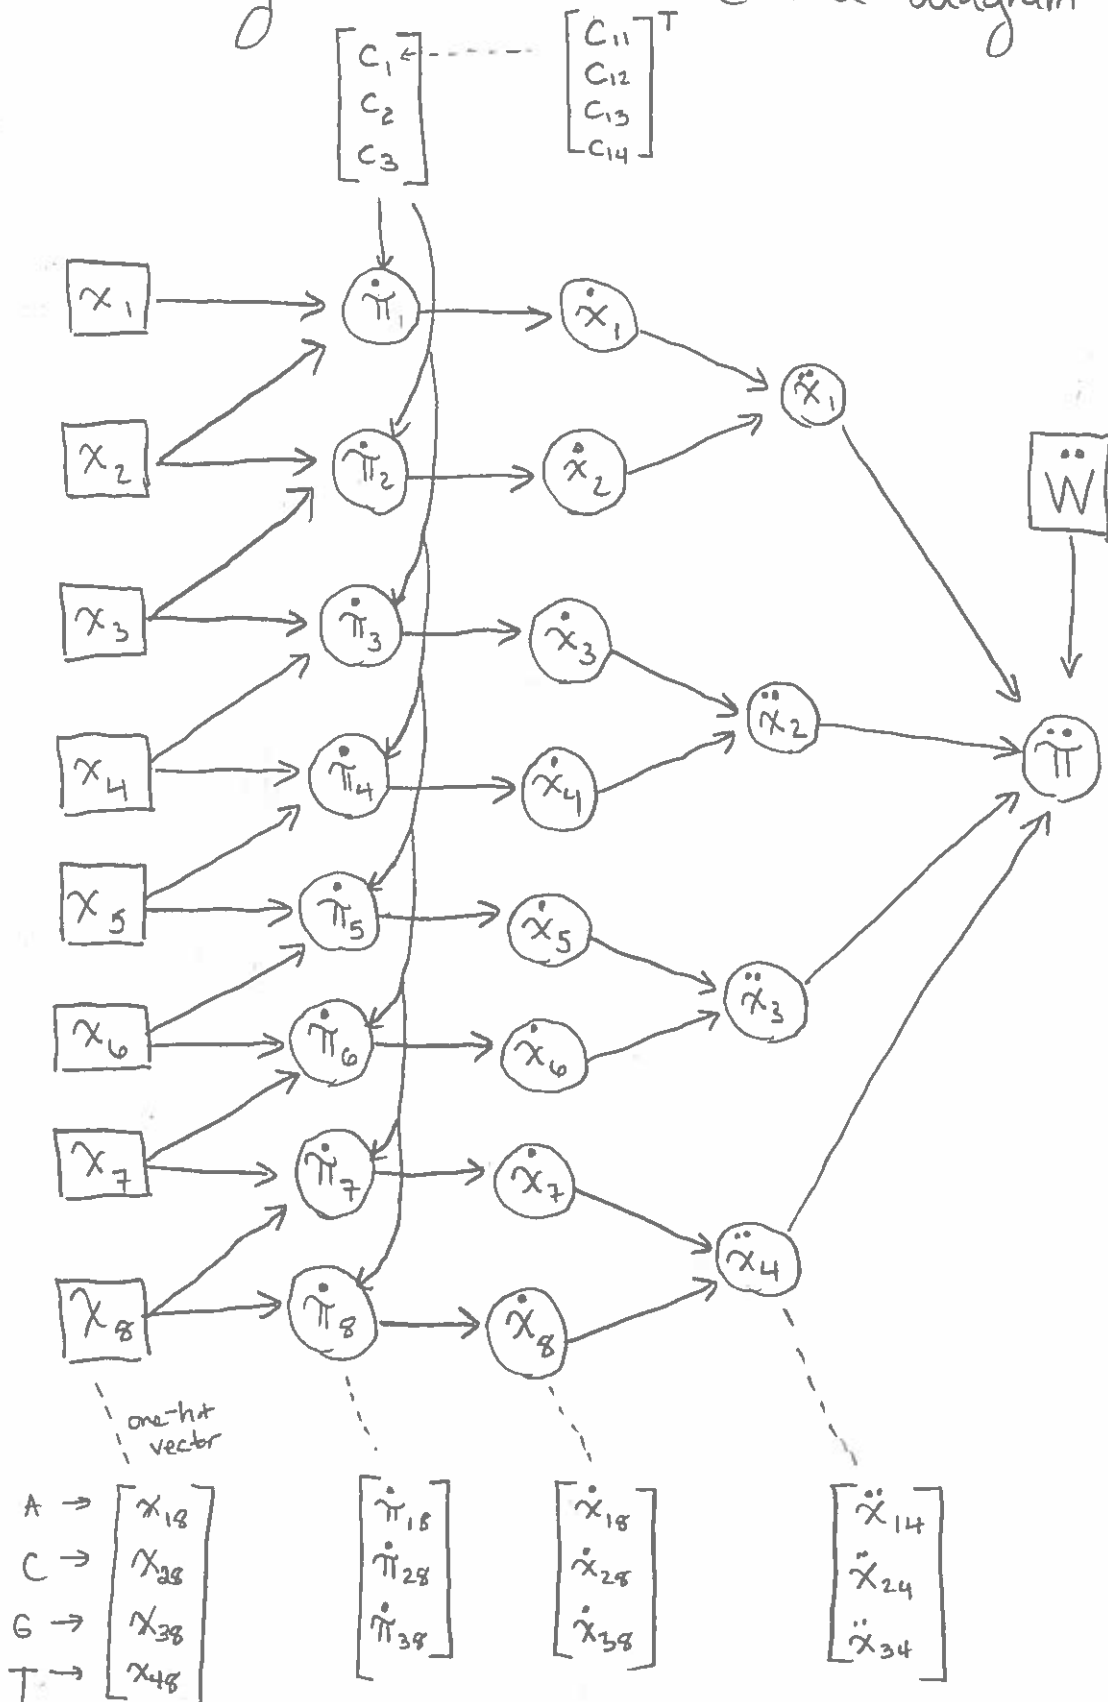
$$- \ddot{x}_{11} = 1 \text{ and } \ddot{x}_{22} = 1 \text{ and } \ddot{x}_{34} = 1$$

$$- \ddot{x}_{11} = 1 \text{ and } \ddot{x}_{23} = 1 \text{ and } \ddot{x}_{34} = 1$$

$$- \ddot{x}_{12} = 1 \text{ and } \ddot{x}_{23} = 1 \text{ and } \ddot{x}_{34} = 1$$


PADDING AND POOLING

⑫ Summarizing all this as a causal diagram gives us:



PADDING AND POOLING

- ⑬ How does one backpropagate through a maxpool layer?
Consider the following simple example:



where $\dot{x} = \max(x_1, x_2)$, and suppose we need to compute $\frac{\partial \dot{x}}{\partial w}$. Assume x_1, x_2 are scalars.

- ⑭ From the Chain Rule:

$$\frac{\partial \dot{x}}{\partial w} = \frac{\partial \dot{x}}{\partial x_1} \frac{\partial x_1}{\partial w} + \frac{\partial \dot{x}}{\partial x_2} \frac{\partial x_2}{\partial w}$$

We know:

$$\dot{x} = \begin{cases} x_1 & \text{if } x_1 \geq x_2 \\ x_2 & \text{if } x_2 \geq x_1 \end{cases}$$

So:

$$\frac{\partial \dot{x}}{\partial x_1} = \begin{cases} 1 & \text{if } x_1 > x_2 \\ 0 & \text{if } x_2 < x_1 \\ \text{undefined} & \text{if } x_1 = x_2 \end{cases} \quad \frac{\partial \dot{x}}{\partial x_2} = \begin{cases} 1 & \text{if } x_2 > x_1 \\ 0 & \text{if } x_1 > x_2 \\ \text{undefined} & \text{if } x_1 = x_2 \end{cases}$$

The functions are piecewise differentiable, much like ReLU:

$$\frac{d}{dz} a(z) = \begin{cases} 1 & \text{if } z > 0 \\ 0 & \text{if } z < 0 \\ \text{undefined} & \text{if } z = 0 \end{cases}$$

PADDING AND POOLING

(15) As with ReLU, we just need to hope that we don't reach a place in our weight space (during gradient descent) where $x_1 = x_2$. Since they're both real numbers, chances are low.

(16) It could be asked:

if we can backpropagate through max, why did we need softmax?



The answer lies in the fact that softmax is really a misnomer for softargmax.

well, softmax is really a misnomer for softargmax...



PADDING AND POOLING

- ⑪ Recall that softmax takes a vector of reals and produces a skewed distribution of the same dimension:

e.g.

$$\begin{bmatrix} 1.6 \\ -0.2 \end{bmatrix} \longrightarrow \begin{bmatrix} .86 \\ .14 \end{bmatrix}$$


This is an approximation of an argmax function that produces a one-hot vector:

$$\begin{bmatrix} 1.6 \\ -0.2 \end{bmatrix} \longrightarrow \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

It is not an approximation of the max function:

$$\begin{bmatrix} 1.6 \\ -0.2 \end{bmatrix} \longrightarrow 1.6$$

- ⑫ Let's try implementing this argmax function:



$$\text{where: } \hat{x}_1 = \begin{cases} 1 & \text{if } x_1 \geq x_2 \\ 0 & \text{if } x_2 > x_1 \end{cases}$$

$$\hat{x}_2 = \begin{cases} 1 & \text{if } x_2 \geq x_1 \\ 0 & \text{if } x_1 > x_2 \end{cases}$$

PADDING AND POOLING

(19) That means:

$$\frac{\partial \dot{x}_1}{\partial x_1} = 0 \quad \frac{\partial \dot{x}_1}{\partial x_2} = 0$$

$$\frac{\partial \dot{x}_2}{\partial x_1} = 0 \quad \frac{\partial \dot{x}_2}{\partial x_2} = 0$$

All the derivatives are degenerate and useless.